


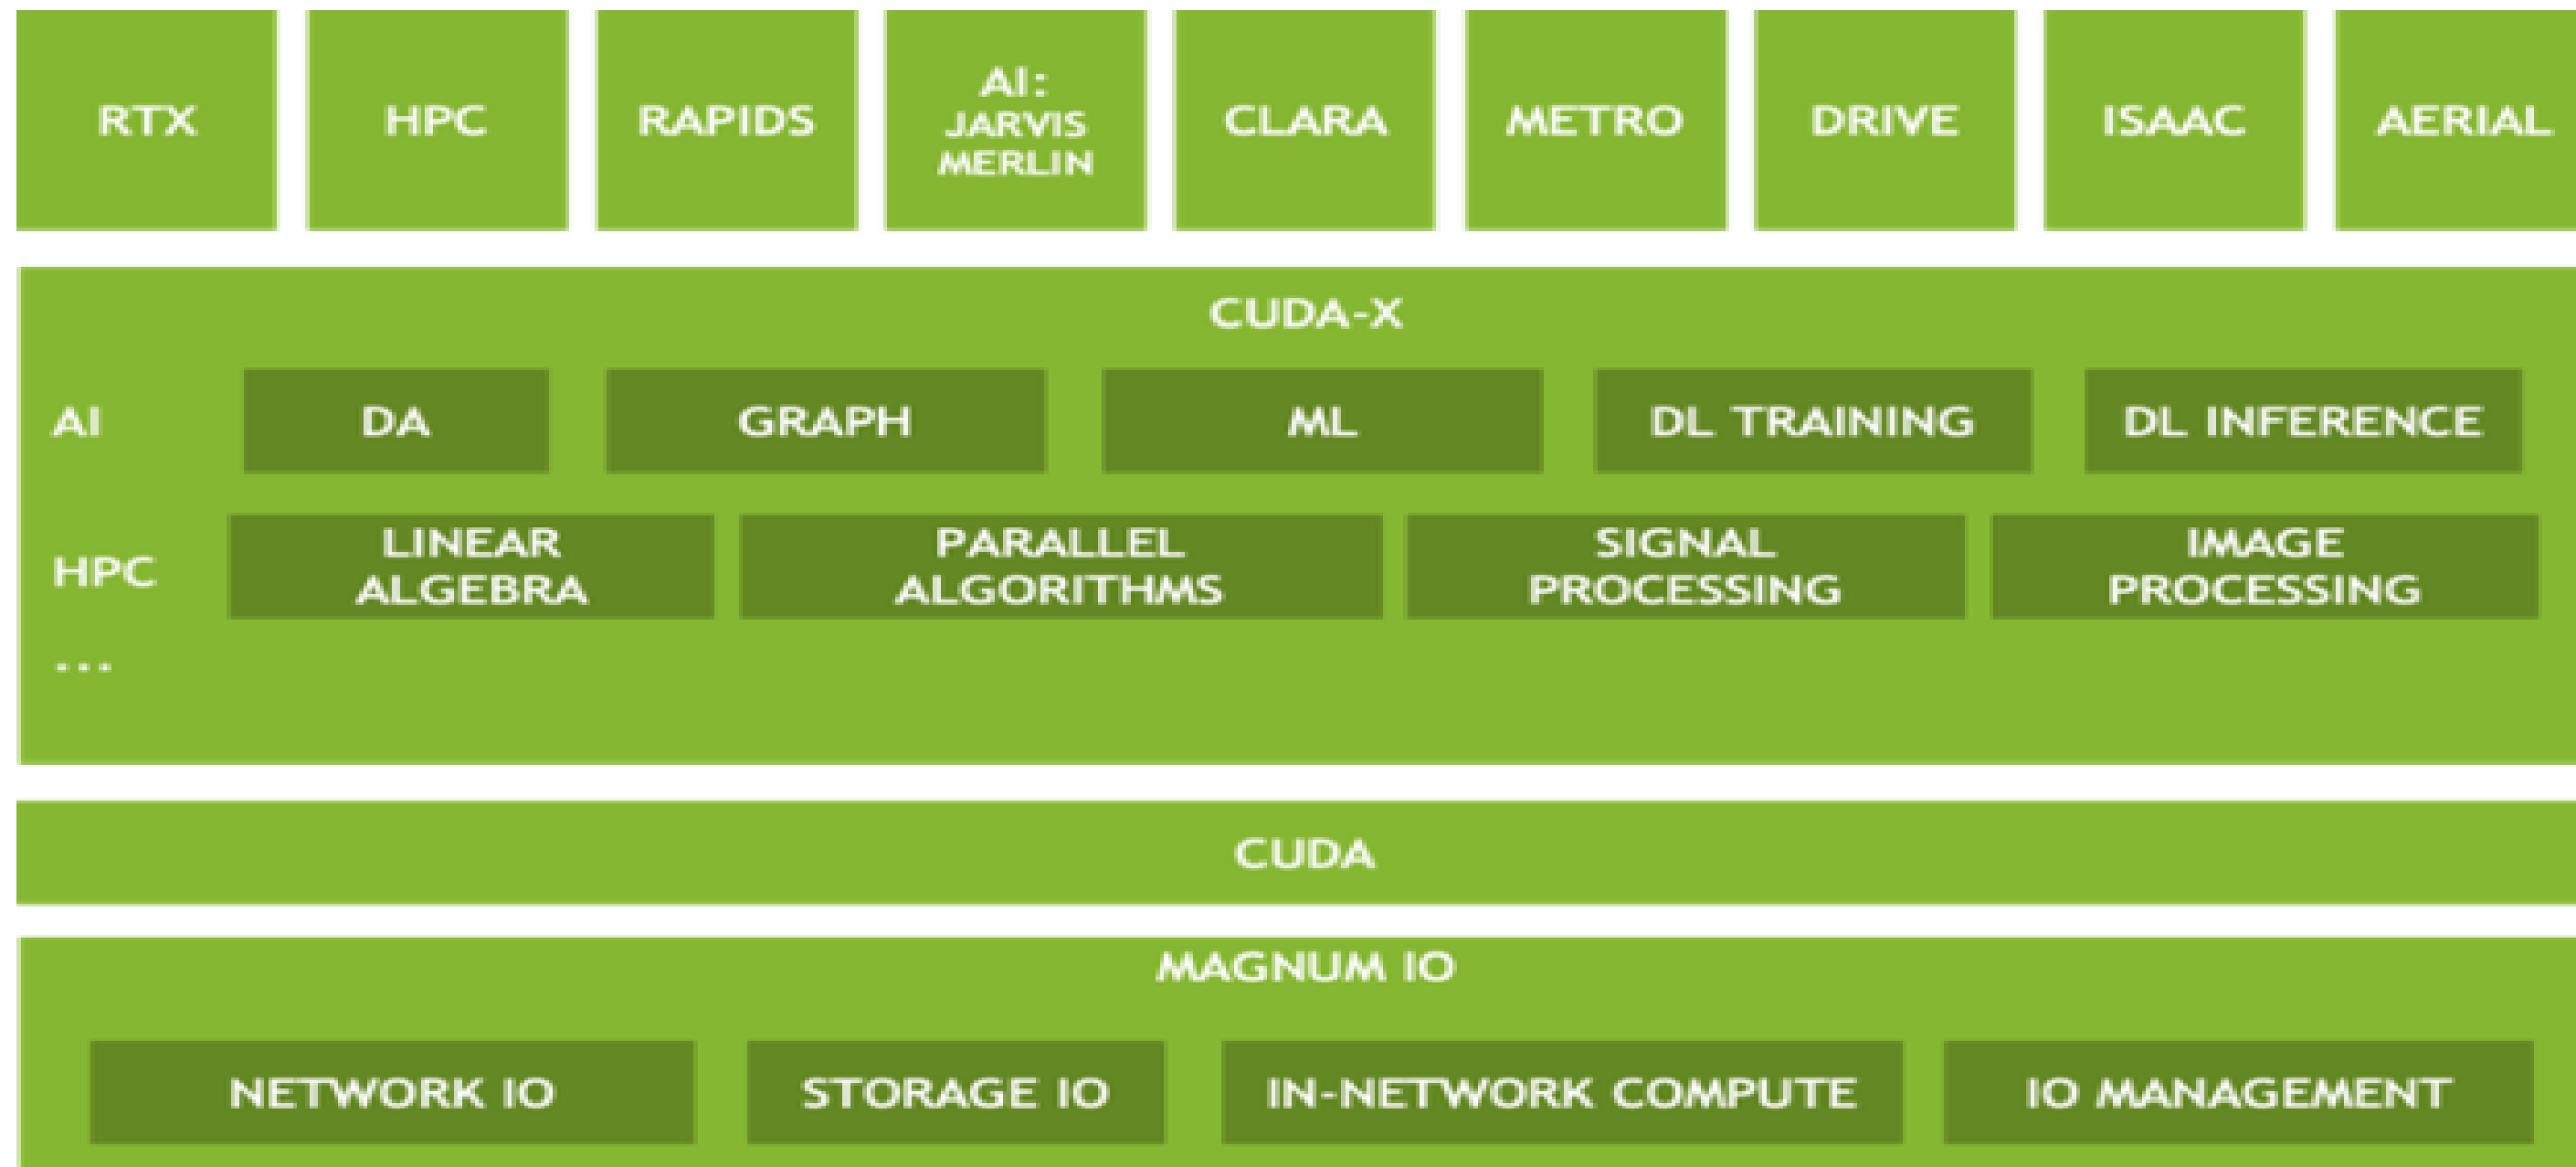


# NVIDIA GDS

Sungta Tsai, Staff Field Application Engineer

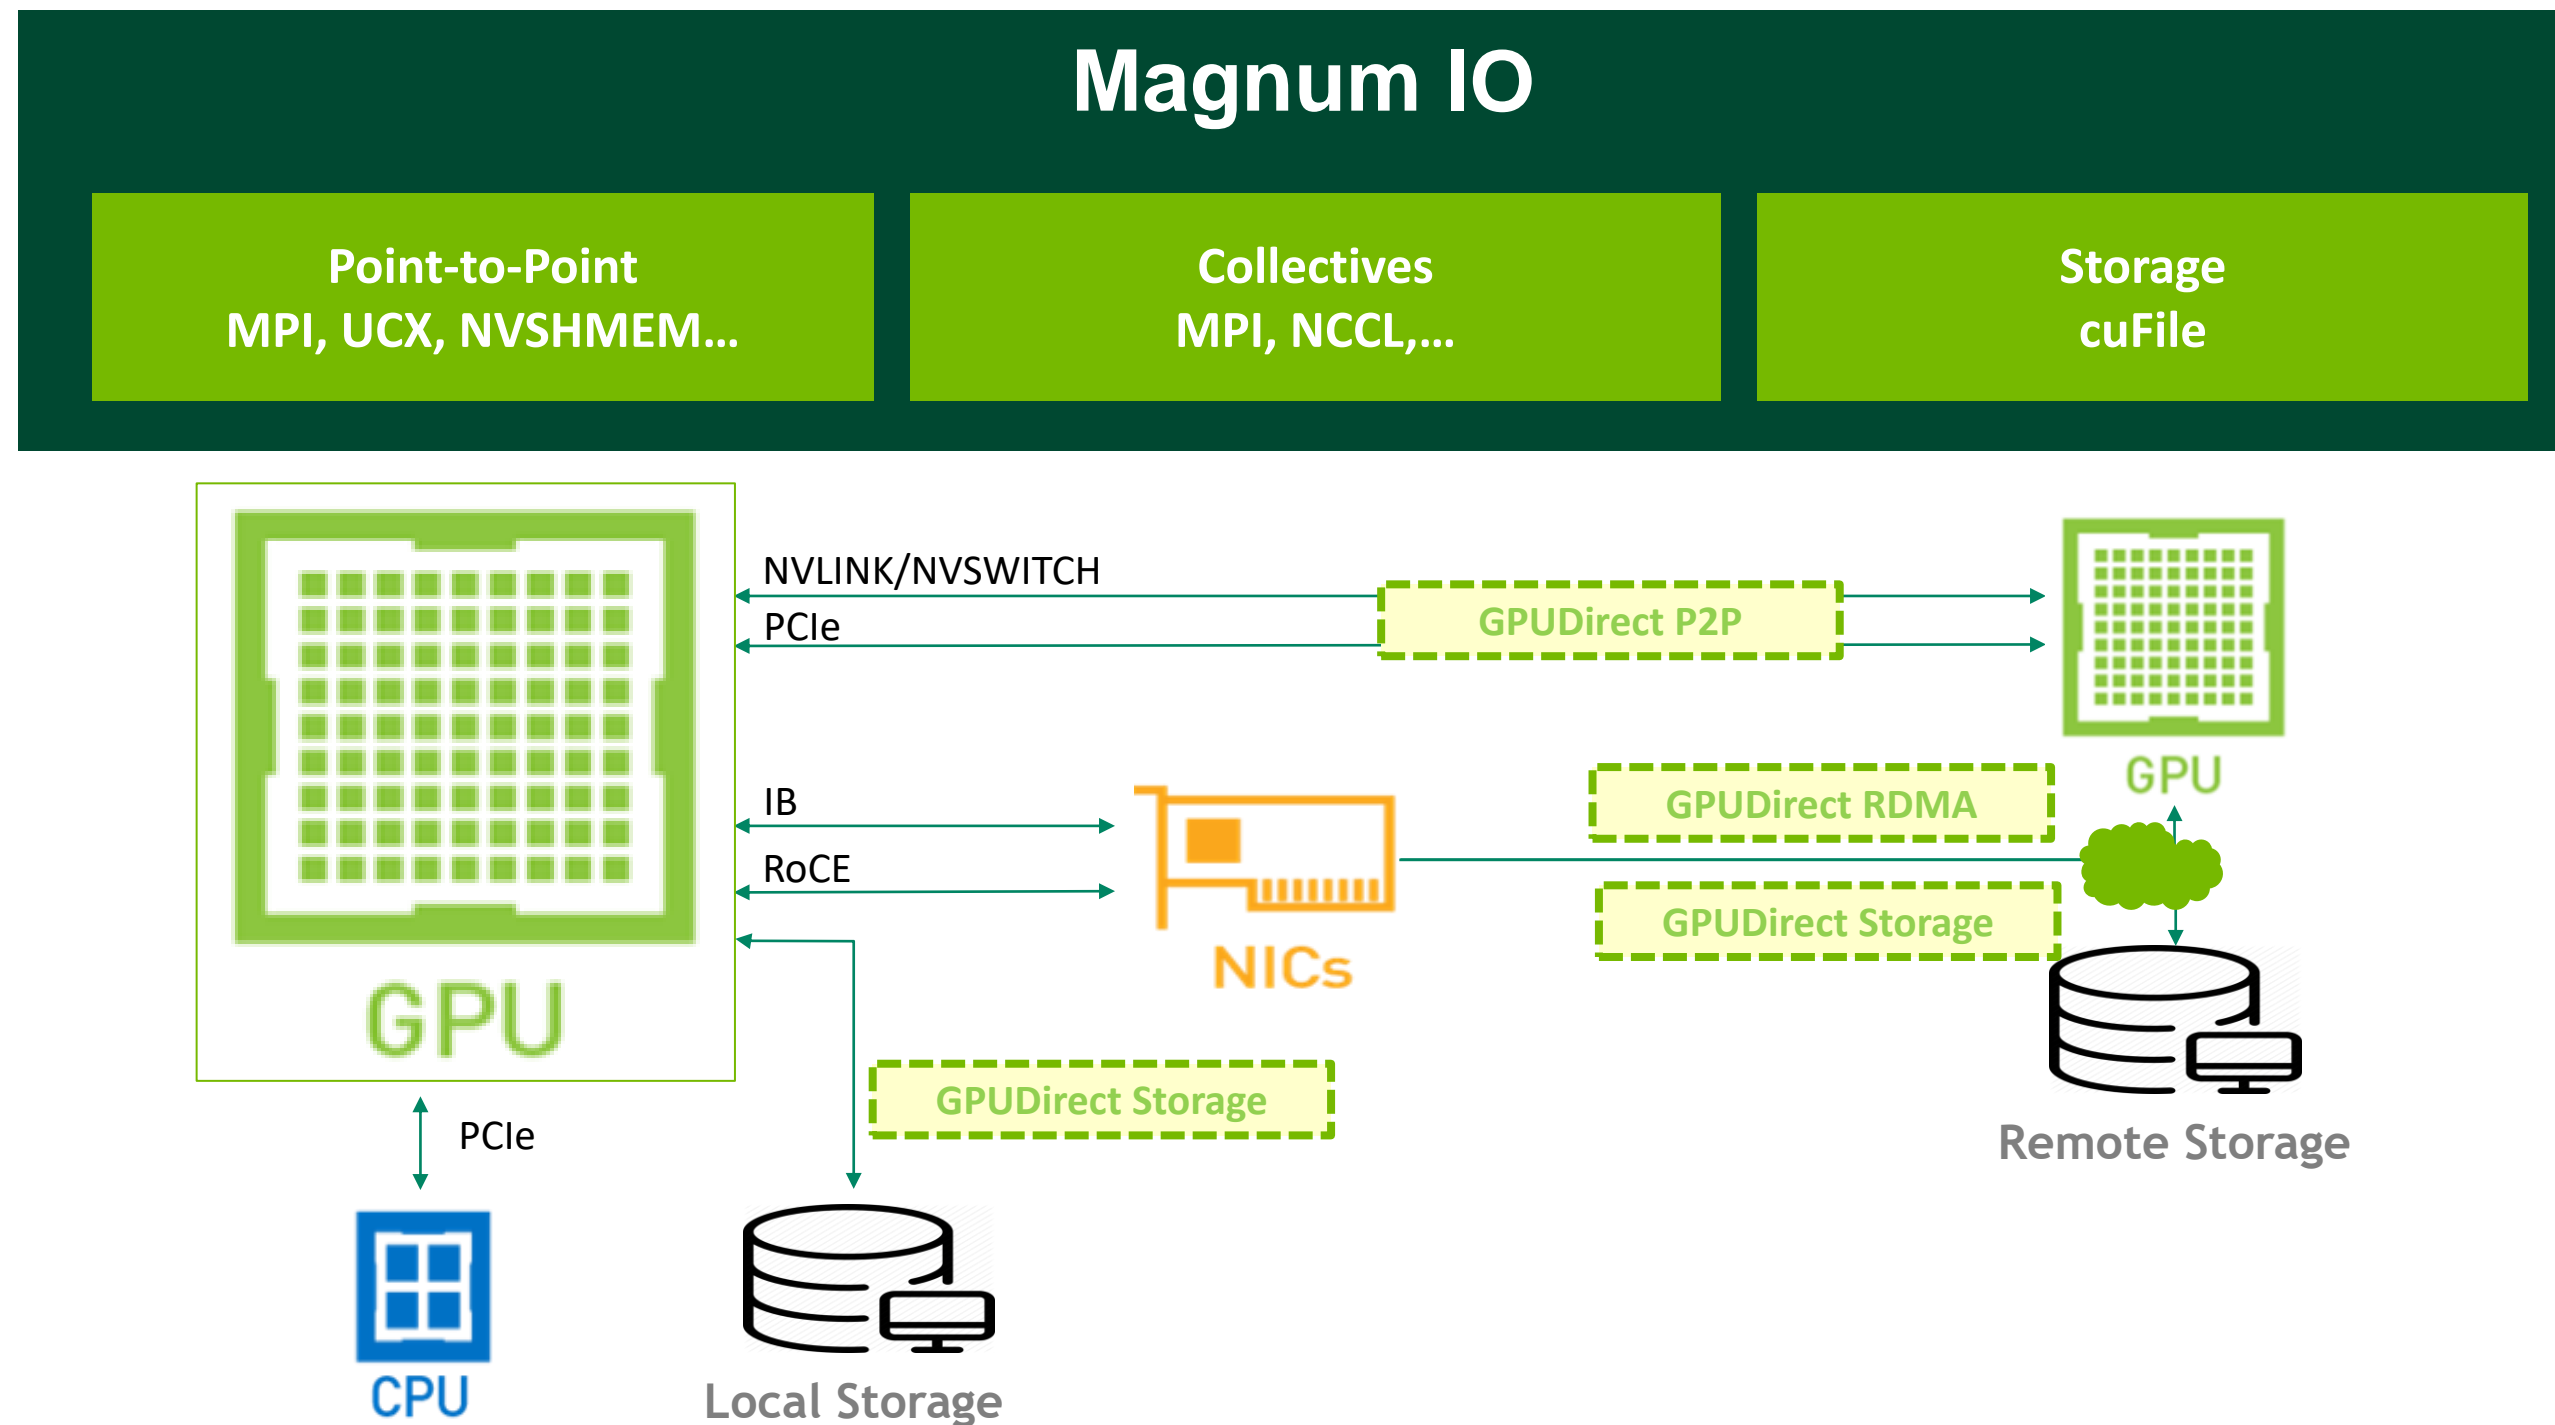


# MAGNUM IO



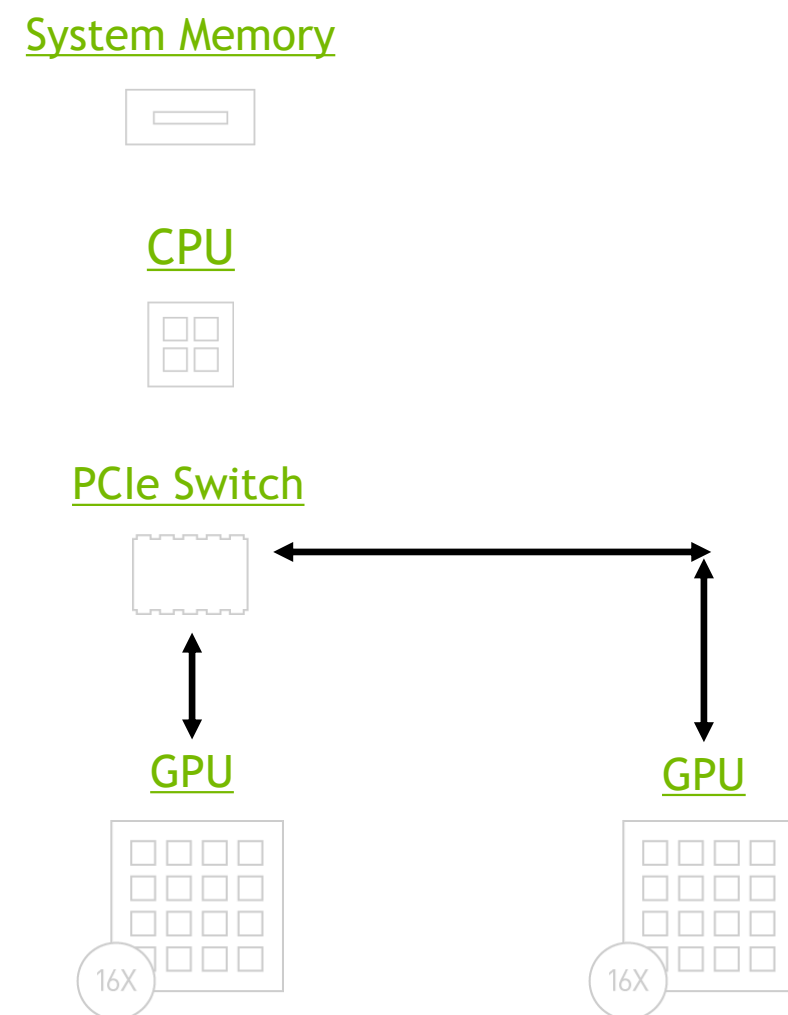
# MAGNUM IO

NVIDIA's Multi-GPU, Multi-Node Networking and Storage IO Optimization Stack

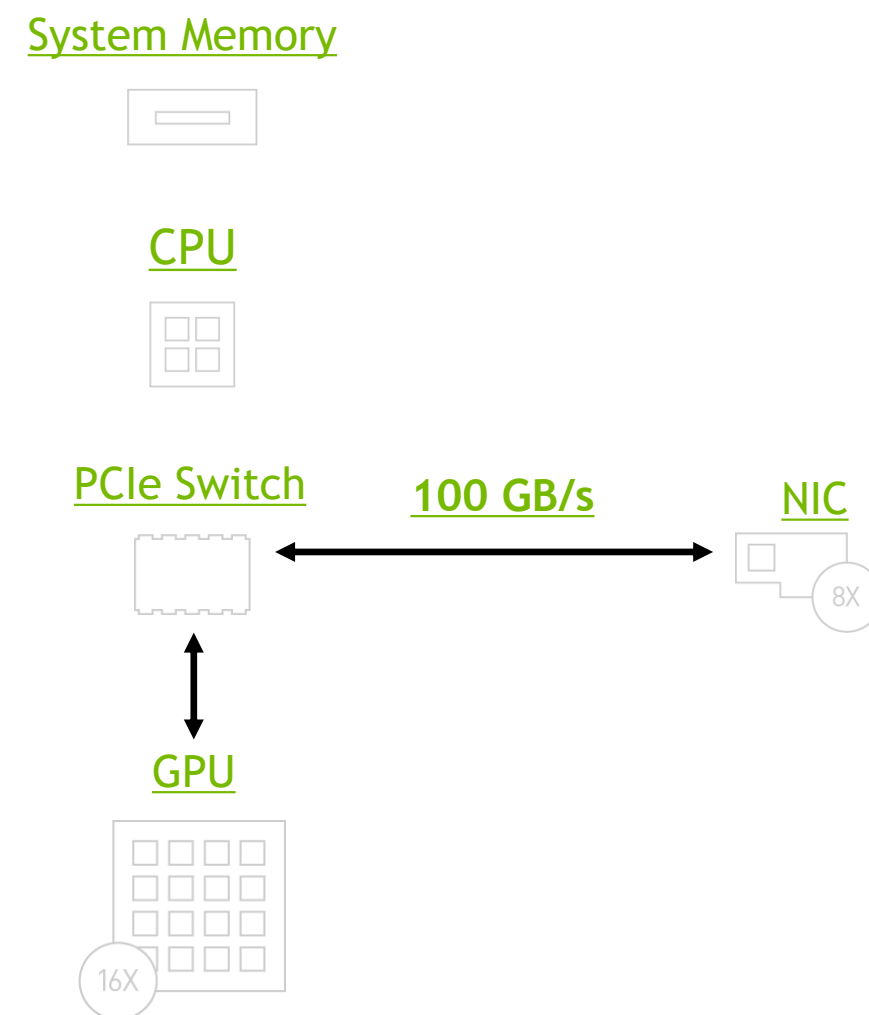


# DATA I/O ACCELERATION

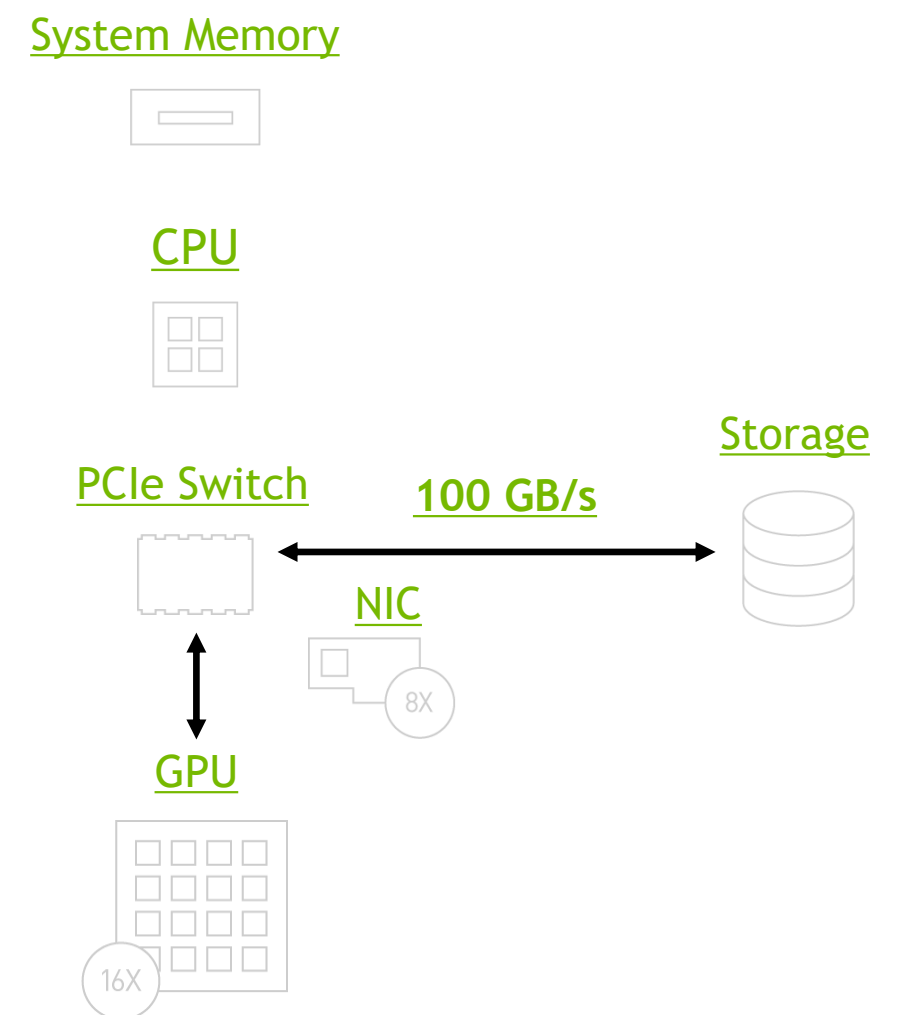
GPU Direct - P2P (Multi-GPU)



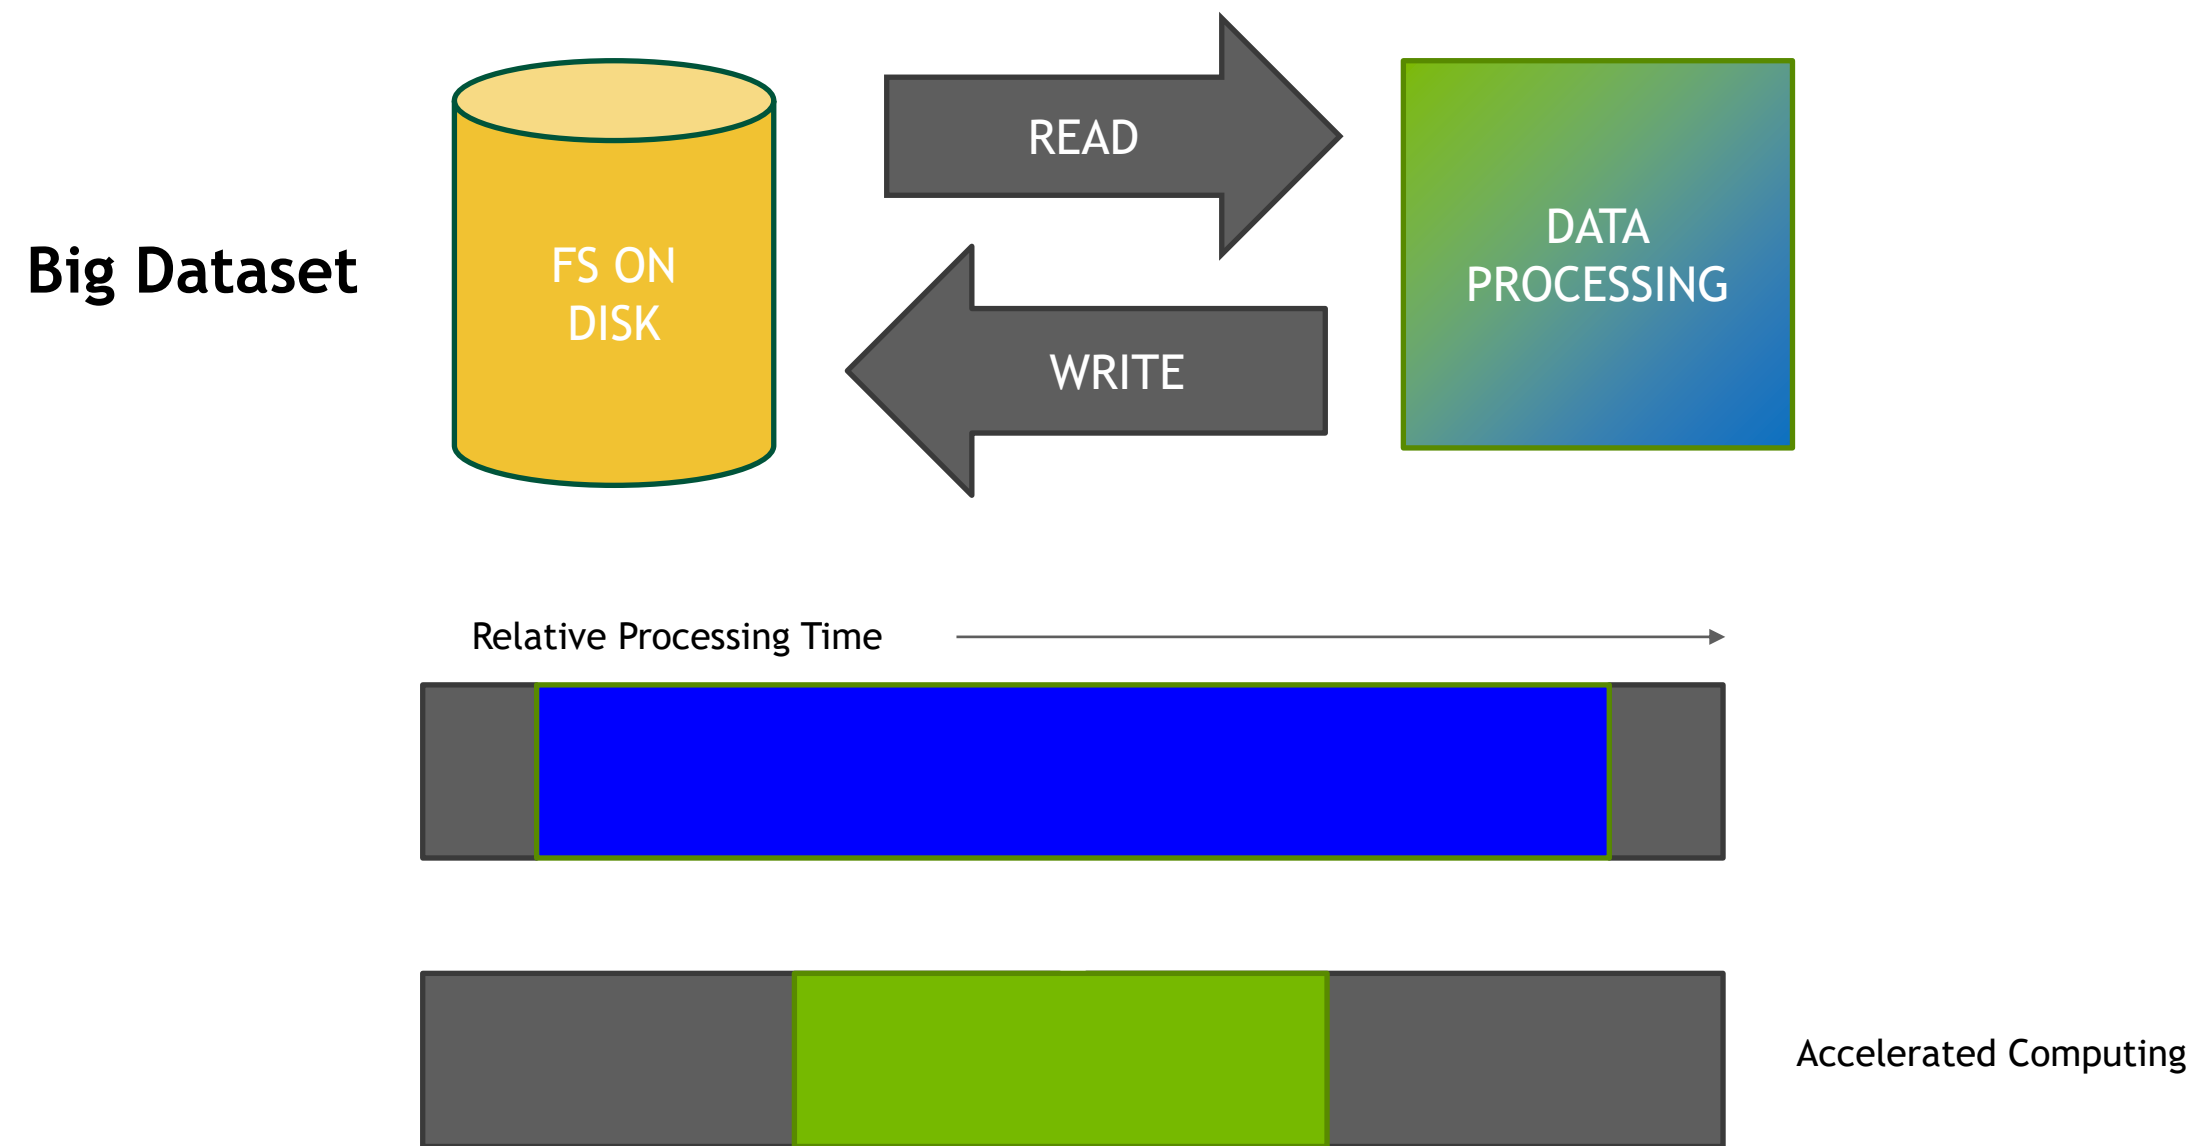
GPU Direct RDMA (Multi-Node)



GPU Direct Storage (Storage)



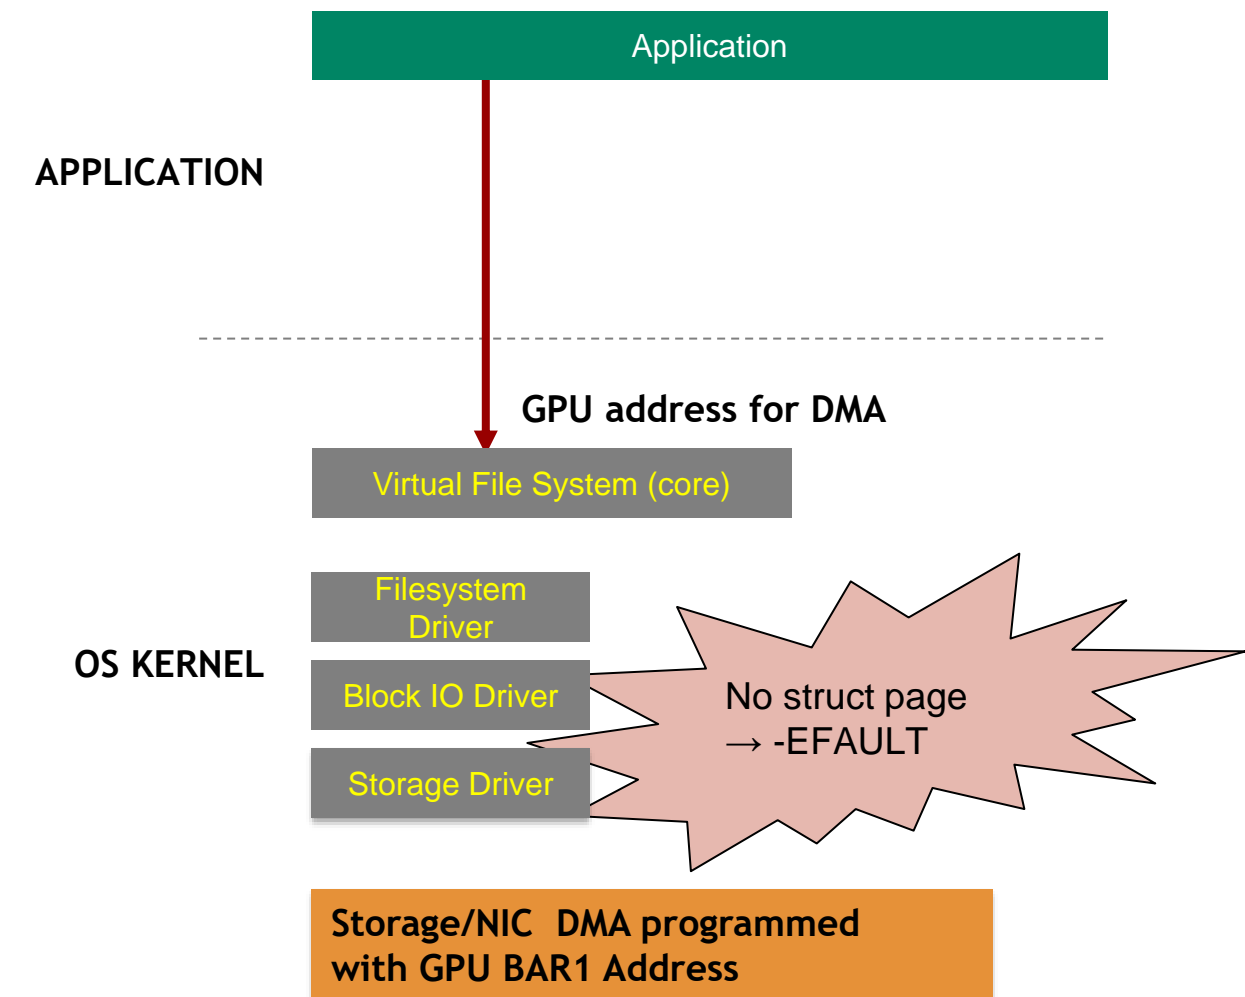
# THE IO CHALLENGE



# GPUDIRECT STORAGE SW ARCHITECTURE

## User and kernel components

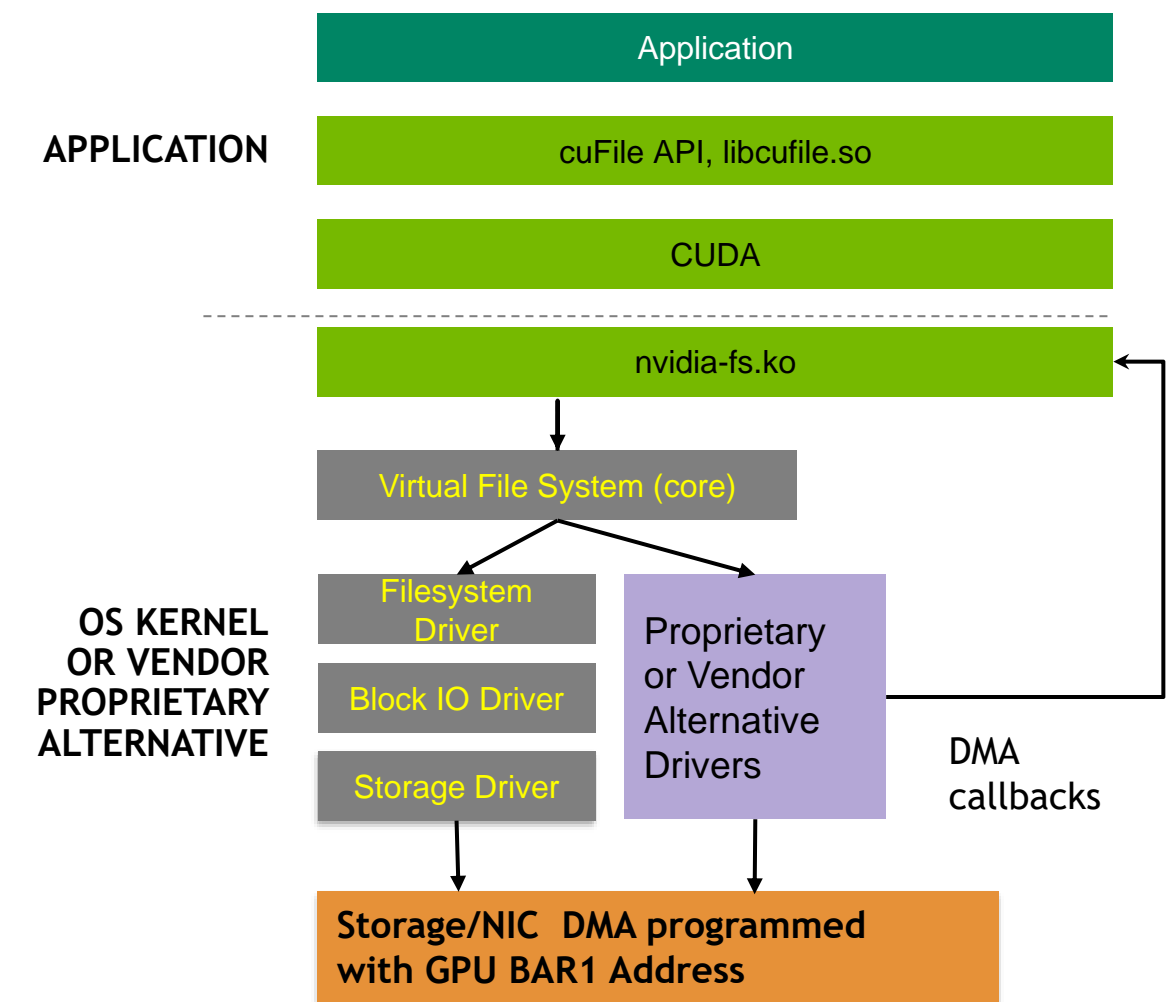

- ▶ Want to program DMA near storage to push/pull data in GPU memory
- ▶ Linux is not enabled to handle GPU Virtual Addresses needed for DMA
- ▶



# GPUDIRECT STORAGE SW ARCHITECTURE

## User and kernel components

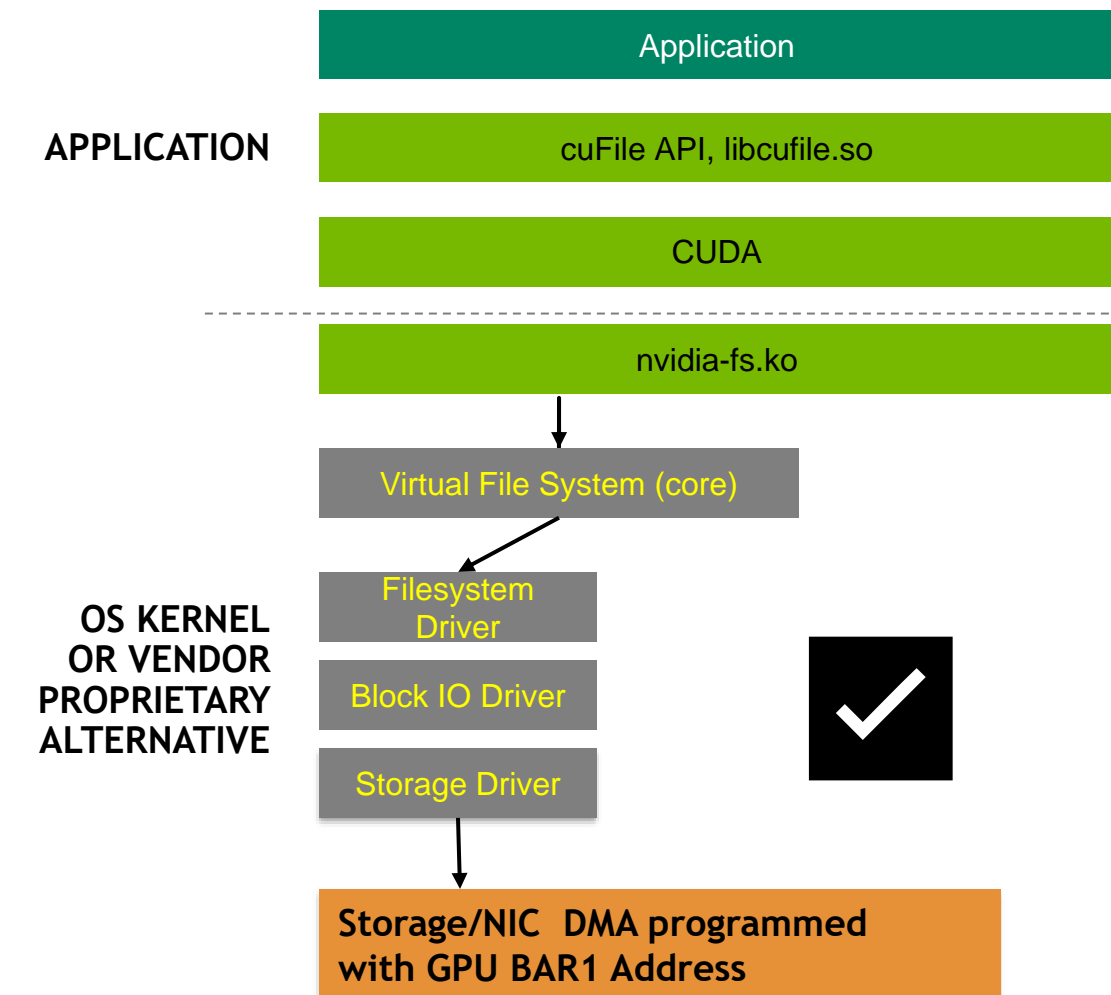

- ▶ cuFile API  
Enduring API for applications and frameworks
- ▶ nvidia-fs Driver API  
For filesystem and block IO drivers  
Vendor-supported solutions: no patching  
avoid lack of Linux enabling
- ▶ NVIDIA is actively working with the community on upstream first to enable Linux to handle GPU VAs for DMA



# GPUDIRECT STORAGE SW ARCHITECTURE

## User and kernel components

- ▶ cuFile interfaces will endure
- ▶ NVIDIA is actively working with the community to enable Linux to handle GPU Virtual Addresses needed for DMA
- ▶ We are increasingly open sourced and are partnering with MLNX in upstreamed efforts






# DATA TRANSFER WITH GDS

## Data Transfer without GPUDirect Storage


- `fd = open("file.txt", O_RDONLY);`
- `buf = malloc(size);`
- `pread(fd, buf, size, 0);`
- `cudaMalloc(d_buf, size);`
- `acudaMemcpy(d_buf, buf, size, cudaMemcpyHostToDevice);`




# DATA TRANSFER WITH GDS

## Data Transfer GPUDirect Storage

- `fd = open("file.txt", O_RDONLY);`
- ~~`buf = malloc(size);`~~ Don't need to have "bounce buffer"
- `pread(fd, buf, size, 0);`
- `cudaMalloc(d_buf, size);`
- `acudaMemcpy(d_buf, buf, size, cudaMemcpyHostToDevice);`



- `fd = open("file.txt", O_RDONLY | O_DIRECT, ...);`
- `cudaMalloc(d_buf, size);`
- `cuFileRead(fhandle, d_buf, size, 0, 0);`



# cuFile APIs

```
/* cuFile File Registration APIs */
int cuFileImportExternalFile(CUFileHandle_t *fh, CUFileDescr_t *descr);
void cuFileDestroyFile(CUFileHandle_t fh);

/* core IO APIs */
ssize_t cuFileRead(CUFileHandle_t fh, void *devPtr, size_t size, off_t offset);
ssize_t cuFileWrite(CUFileHandle_t fh, void *devPtr, size_t size, off_t offset);

/* APIs to Register user specified buffer for direct BAR1 mapping */
CUfileError_t cuFileBufRegister(void *devPtr, size_t size, int flags);
CUfileError_t cuFileBufDeregister(void *devPtr);

/*APIs to control the Driver and cuFile resource lifecycle */
CUfileError_t cuFileDriverOpen();
void cuFileDriverClose();

/*APIs to tune the Driver and cuFile resource usage*/
CUfileError_t cuFileGetDriverProperties(CUfileDrvProps_t *props);
CUfileError_t cuFileDriverSetPollMode(bool poll, size_t poll_threshold_size);
CUfileError_t cuFileDriverSetMaxDirectIOSize(size_t max_direct_io_size);
CUfileError_t cuFileDriverSetMaxCacheSize(size_t max_cache_size);
CUfileError_t cuFileDriverSetBAR1Size(size_t max_bar1_size);
```

# USAGE EXAMPLE


Write from GPU memory to local or remote storage

```
int main(void) {
    ...
    CUfileError_t status;
    CUfileDescr_t cf_descr;           // general enough to support Windows
    CUfileHandle_t cf_handle;

    status = cuFileDriverOpen();
    fd = open(TESTFILE, O_WRONLY|O_DIRECT, 0644);           // interop with normal file IO
    cf_descr.handle.fd = fd;
    status = cuFileImportExternalFile(&cf_handle, &cf_descr);
    cuda_result = cuMemAlloc(&devPtr, size);
    status = cuFileBufRegister(devPtr, size);               // performance optimization
    assert(cudaMemset((void *) devPtr, 0xab, size) == cudaSuccess);
    ret = cuFileWrite(cf_handle, devPtr, size, 0);           // ~pwrite: file handle, GPU address, size, offset


    status = cuFileBufDeregister(devPtr);                   // optional cleanup for good hygiene
    cuFree(devPtr);
    cuFileDestroyFile(cf_handle);
    close(fd);
    cuFileDriverClose();
    return 0;
}
```

# GPUDIRECT STORAGE - USE CASES




## NASA Mars Lander

Simulation → visualization  
128TB data, must stream in from remote  
Part render, part IO; not quite linear in IO  
5 GB/s (1 DGX-2) → 160-180 GB/s (4 DGX-2s) with GPUDirect Storage



## Molecular Dynamics

Simulation → analytics → visualization  
30TB data, can be remote + local  
 $O(N^2)$  IO problem to build dissimilarity matrix of macromolecule poses across time steps so can find stable configs  
7 GB/s → 22 GB/s local, 60 GB/s remote  
3x from GDS vs. heroic effort, 4x threads



## Pangeo Earth Science

Simulation → DA, DL → visualization  
100 TB-PB data, streamed from remote  
Coming to GPUs because of faster IO  
Increasing richness: DA, DL  
Moving from 1 per day to 2-3 per day  
What ifs vs. safe bets

# IO500

IO to compute dominates


DGX-2: 16 GPUs, 8 NICs, 2 CPUs

Move data directly to GPUs

Relieve CPU bottleneck

6 or 10 DGX-2,  
10 DDN EXA5 on A3I AI400X  
GPUDirect Storage used for IOR easy

CPU 10 Nodes vs. GDS 6 and 10 Nodes



IOR GDS easy read and easy write scaling





# TPC-H (REAL BUT EXTREME CASE)

Speedups from both IO and savings in CPU memory management

## Q4 TPC-H Benchmark Work Breakdown: With Repeated Query

